Finding Perfect Squares using Elixir

COP5615 – Distributed Operating Systems Principles Project 1

The goal of this project is to use Elixir and the actor model to build a good solution to find sequence of squares of integers that leads to a perfect square and runs well on multi-core machines.

For the entire problem description please read problemStatement1.pdf

Authors

- **Akash R Vasishta** – *UF ID: 5395 5080*

- **Aditya Bhanje** – *UF ID: 9697 1842*

Algorithm

In this program, after taking input from the user in the form of (N,k), where k is the number of integers in the sequence and N is the final number for the start of the sequence, we designate each actor process a range of sequences(work unit) for which it has to find whether the sum of the squares of those integers is a square of an integer or not. If the result comes out to be true, it then prints the start number of that sequence. Once all the sequences given to that actor are computed, it reports the process which spawned it about it's exit.

Size of the work unit and Running Time

For the problem of **N=1000000 and k=24**, the following table shows the relation between Work Unit,CPU TIME(SYS TIME + USER TIME), REAL TIME and CPU UTILISATION

Work Unit	CPU TIME	REAL TIME	CPU UTILISATION
10	31.7	29.52	1.07
100	1.4	0.63	2.22
1000	1.115	0.492	2.26
10000	1.124	0.503	2.23
50000	1.088	0.481	2.26
100000	1.097	0.506	2.16

For the problem of **N=10000000 and k=24**, the following table shows the relation

between WORKUNIT, CPU TIME(SYS TIME + USER TIME), REAL TIME and CPU UTILISATION

Work Unit	CPU TIME	REAL TIME	CPU UTILISATION
100	44.16	34.91	1.26
1000	10.73	2.96	3.625
10000	10.6	2.92	3.63
20000	10.61	2.93	3.62
100000	10.59	2.92	3.62

After observing these results, we came to the conclusion of choosing the size of work unit to be 20000, as for larger values of N it gave us a good CPU Utilisation value.

For work unit as **20000** and **k=4**, the following was observed:

N	CPU TIME	REAL TIME	CPU UTILISATION
10	0.444	0.324	1.37
100	0.451	0.329	1.37
1000	0.434	0.327	1.32
10000	0.451	0.328	1.37
100000	0.502	0.348	1.44
1000000	1.104	0.492	2.24
10000000	10.606	2.936	3.612
100000000	108.784	28.23	3.853
1000000000	1145.937	296.038	3.87

Largest number of working machines on which the code was run

The code was simultaneously run on 2 machines with the actors equally divided among the 2 machines.

Output of N = 1000000, k = 4

rvakash@rvakash-HP-Pavilion-15-Notebook-PC:~/Documents/masters/Masters/UF/DOS/projects/Find PerfectSquares/proj1$ time mix run -e Proj1.scheduler 1000000 4 0

real   0m0.623s

user   0m1.190s

sys   0m0.059s

rvakash@rvakash-HP-Pavilion-15-Notebook-PC:~/Documents/masters/Masters/UF/DOS/projects/Find PerfectSquares/proj1$

Largest Problem successfully solved

The above code successfully runs for the values of **N=1000000000, k=24** and CPU Utilisation was found out to be **3.87** for the same.

Running the code

After unzipping the provided folder, open the command prompt in that folder and type the following commands:

$ mix run lib/proj1.ex <arg1> <arg2> <arg3>

where arg1 is the value of N and arg2 is the value of k.

arg3– 0 for local, 1 for multiple machines.

Example – time mix run lib/proj1.ex 1000000 24 0

Running on Multiple machines(demo will be given during TA Hours)

1) Machine – 1 and Machine –2 should be on the same network
2) Machine – 1

        a. $iex –name node_one@192.168.1.2 –cookie something
3) Machine – 2
        a. $ iex –name node_two@192.168.1.47 –cookie something
4) Machine – 1
        a. $Node.connect :node_two@192.168.1.47
5) Compile the code on both the machines –
        a. $c("perfectSquares.exs")
6) Run on Machine – 1
        a. $PerfectSquares.scheduler(1000000, 24, 1)

Where node_one@192.168.1.2 and node_two@192.168.1.47 are the IP Addresses of the two machines respectively.